

# Lab 4 Trainer Demo and Answer Key

Private facilitation guide for the tone-of-voice lab.

**TRAINER ONLY**

Do not distribute this file with the participant pack. Use the expected findings to facilitate discussion, not to force one perfect writing style.

| Timing | Activity                   |
|--------|----------------------------|
| 5 min  | Introduction and baseline  |
| 10 min | Sample analysis            |
| 8 min  | Human rule validation      |
| 8 min  | Profile and Notebook setup |
| 10 min | Apply and score            |
| 10 min | Controlled update          |
| 7 min  | Audience test and debrief  |

## Demonstration opening

**Say:**

"Today we are not asking Copilot to sound vaguely like us. We will give it evidence, decide which patterns are intentional, test the result and then update one rule without changing the underlying facts."

**Ask:**

- What is the risk of learning a voice from one email?
- What is the difference between voice and tone?
- How would you know that an updated tone is actually better?

## Expected baseline weaknesses

| Weakness             | Expected explanation                                                      |
|----------------------|---------------------------------------------------------------------------|
| Formal opening       | Copilot may use a generic business-email pattern.                         |
| Too much reassurance | It may convert "does not currently affect" into a stronger guarantee.     |
| Unnecessary apology  | A delay can trigger apologetic language even when no apology is required. |
| Weak next step       | The owner and confirmation deadline may be buried.                        |
| Generic closing      | The message may not sound like an established client relationship.        |

**DO NOT GRADE EXACT WORDING**  
The objective is evidence, control and usability. Several different drafts may be valid.

## Expected candidate tone rules

| Rule                                       | Examples                      | Decision                       |
|--------------------------------------------|-------------------------------|--------------------------------|
| Lead with the practical update             | Samples 1, 3 and 8            | Accept                         |
| Give enough context to understand impact   | Samples 1, 2 and 4            | Accept                         |
| Separate facts from recommendations        | Samples 4, 5 and 8            | Accept                         |
| State owners and deadlines only when known | Samples 1, 3 and 8            | Accept                         |
| Always use bullets                         | Not consistent across samples | Reject                         |
| Always avoid technical terminology         | Sample 2 uses necessary terms | Edit                           |
| Use a warm but short closing               | Samples 1-4                   | Accept with audience condition |
| Copy sample openings exactly               | Case-specific wording         | Reject                         |

# Recommended Tone Profile v1.0

Approved instruction

Write in Hanna's professional consulting voice. The voice is friendly, human and confident, but not overly formal or promotional. Lead with the outcome, status or reason for the message. Use clear international English and concise sentences. Provide enough context to prevent misunderstanding. When communicating a problem, state what changed, explain the practical impact and identify the next action, owner and deadline when known. Separate confirmed facts, recommendations and unresolved questions. Do not invent certainty, commitments, owners, deadlines or approvals. Avoid corporate filler, excessive apologies, exaggerated enthusiasm and generic AI expressions. Use bullets only when several actions or decisions are easier to scan. End with a clear, natural next step.

| Factor     | Recommended rating | Notes                                           |
|------------|--------------------|-------------------------------------------------|
| Warmth     | 3/5                | Human and approachable without over-familiarity |
| Directness | 4/5                | Outcome and action are explicit                 |
| Formality  | 3/5                | Professional but not stiff                      |

# Model v1.0 output

EXAMPLE

Hi Alex, A quick update on Project Atlas UAT. The environment will be available on Tuesday morning rather than Monday because the security configuration will finish on Monday afternoon. This does not currently affect the final launch date. Mark will confirm environment readiness by 15:00 on Monday, and I will share the confirmation once it is available. The training-video branding remains outside this readiness decision. Please let me know if you would like us to cover anything else in the readiness update. Best, Hanna

Verification:

- All source facts are present.
- No final security result is invented.
- The launch date is not guaranteed.
- Mark remains the named owner.
- The excluded training topic is not converted into an action.

## Controlled update: v1.0 to v1.1

| Field         | Expected outcome                                                                                                          |
|---------------|---------------------------------------------------------------------------------------------------------------------------|
| Feedback      | Opening and closing feel slightly formal for an existing client.                                                          |
| Rule affected | Opening and closing warmth for established relationships.                                                                 |
| Unchanged     | Facts, length, directness, structure, owner, deadline and purpose.                                                        |
| Current rule  | Begin with a concise explanation of the message purpose.                                                                  |
| Updated rule  | For existing clients, begin with a short natural update and close with a brief collaborative invitation when appropriate. |
| Risk          | Too much warmth may sound casual or add unnecessary sentences.                                                            |
| Test          | Repeat the original Project Atlas email and compare only opening and closing.                                             |

**APPROVAL LOGIC**  
Approve v1.1 only if the message is warmer and all other evaluated dimensions remain equal or improve. Otherwise revise or roll back.

## Model v1.1 output

**EXAMPLE**  
Hi Alex, Just a quick update on Project Atlas UAT. The environment will be available on Tuesday morning rather than Monday because the security configuration will finish on Monday afternoon. This does not currently affect the final launch date. Mark will confirm environment readiness by 15:00 on Monday, and I will share the confirmation once it is available. The training-video branding remains outside this readiness decision. Let me know if there is anything else you would like included in the readiness update. Best, Hanna

**Expected changed sentences:**

- "A quick update" becomes "Just a quick update".
- The closing becomes a natural invitation appropriate to an existing client.
- No other source fact or action changes.

## Scorecard interpretation

| Signal                                           | Facilitator response                                                  |
|--------------------------------------------------|-----------------------------------------------------------------------|
| v1.1 warmth rises; all other scores stable       | Approve the controlled change.                                        |
| v1.1 becomes longer or less direct               | Revise; the change affected more than one parameter.                  |
| Copilot adds an apology or guarantee             | Reject the output; factual and instruction controls failed.           |
| Participants prefer v1.0                         | Accept the evidence and roll back. Newer is not automatically better. |
| Different groups produce different valid wording | Compare outcomes and controls rather than exact phrases.              |

### Facilitation question:

"If a participant likes the new version but cannot explain what changed, is the update controlled?" Expected answer: No. Preference alone does not prove which instruction caused the improvement.

## Audience stress-test expectations

| Stakeholder       | Expected adjustment                                             | Expected stress-test outcome                  |
|-------------------|-----------------------------------------------------------------|-----------------------------------------------|
| Existing client   | Slightly warmer relationship language                           | Clear, calm, outcome-first                    |
| Technical lead    | More detail about security configuration and readiness evidence | Direct, evidence-aware, no invented certainty |
| Executive sponsor | Shorter focus on impact, launch risk and confirmation point     | Confident, concise, decision-oriented         |

# Common problems and recovery

| Problem                                   | Recovery                                                  |
|-------------------------------------------|-----------------------------------------------------------|
| Copilot learns errors as style            | Explicitly exclude typos and case-specific wording.       |
| Every message sounds identical            | Separate stable voice from audience and channel overlays. |
| Participant changes many adjectives       | Return to the fixed case and one-variable prompt.         |
| Notebook instruction is too long          | Keep stable rules; move task details into the prompt.     |
| Custom instructions affect unrelated work | Keep project-specific guidance in the Notebook.           |
| New version is worse                      | Use the version log and restore the previous instruction. |

## Debrief answers

- Voice is the relatively stable communication identity; tone adapts to audience and situation.
- Multiple samples reduce the risk of treating one context as a permanent rule.
- The same source case is necessary for a fair before-and-after comparison.
- One-parameter testing makes the causal change explainable.
- Human approval remains necessary because recognisability and appropriateness are judgement questions.

# Trainer completion checklist

- Participants saved a baseline before viewing the samples.
- At least one candidate rule was edited or rejected.
- The profile includes factual safeguards.
- v1.1 contains exactly one approved rule change.
- The scorecard compares baseline, v1.0 and v1.1.
- Participants can explain when to use Notebook versus personal custom instructions.
- The group discussed a rollback scenario.

### CLOSE WITH

"A useful tone profile is not a static description of personality. It is a tested set of communication rules that you can explain, update and reverse."